

CS144 2020

Reliability **from unreliability**

Keith & Nick

What the Internet provides

- “Best-effort” delivery of datagrams
 - up to about 1,500 bytes
 - from one computer to another
- “Best-effort” means datagram **might** be:
 - **lost completely**
 - delivered **more than once**
 - **delivered *after* another datagram that was sent later**
 - delivered with some **bytes changed**
 - delivered but **truncated**

What most users and applications want

- Reliable retrieval of a short piece of data
 - “What’s the IP address that corresponds to cs144.keithw.org?”
- Reliable action
 - The text of Keith’s message #7 is : “Fire a torpedo!”
- **Reliable byte stream**
 - Sequence of bytes (in each direction) delivered in order, correctly
- Reliable delivery of a large file (FTP, SMTP, HTTP)
 - “Subject: Homework. Dear Professor McKeown: Here is my 20 MB file.”
 - “Subject: Re: Homework. Thank you! Please call me Nick.”
- Reliable remote procedure call (RPC) (HTTP/1, HTTP/2, HTTP/3, gRPC, Thrift)
 - `POST /bankaccounts/checking/billpay HTTP/1.1`
`amount=270,000&payee=StanfordSailing&memo=admitmychildplz`

Reliability

- A module behaves **reliably** when it:
 - provides **some** stated abstraction/interface
 - even in the face of underlying faults (e.g. packet loss)
 - and when it can't do that, the module signals failure.

The big question

How to provide these abstractions
reliably on top of an **unreliable** system?

TCP in a nutshell

» datagram that says ~~bytes 0..49 of the byte stream have the contents: “MAIL FROM: <thepope@vatican.va>”~~

» datagram that says bytes 50..99 of the byte stream have the contents: “DATA\nHi Keith here is your ordination.”

« “The next byte of the stream that I need from you is #0.”

» datagram that says bytes 0..49 of the byte stream have the contents: “MAIL FROM: <thepope@vatican.va>”

« “The next byte of the stream that I need from you is #100.”